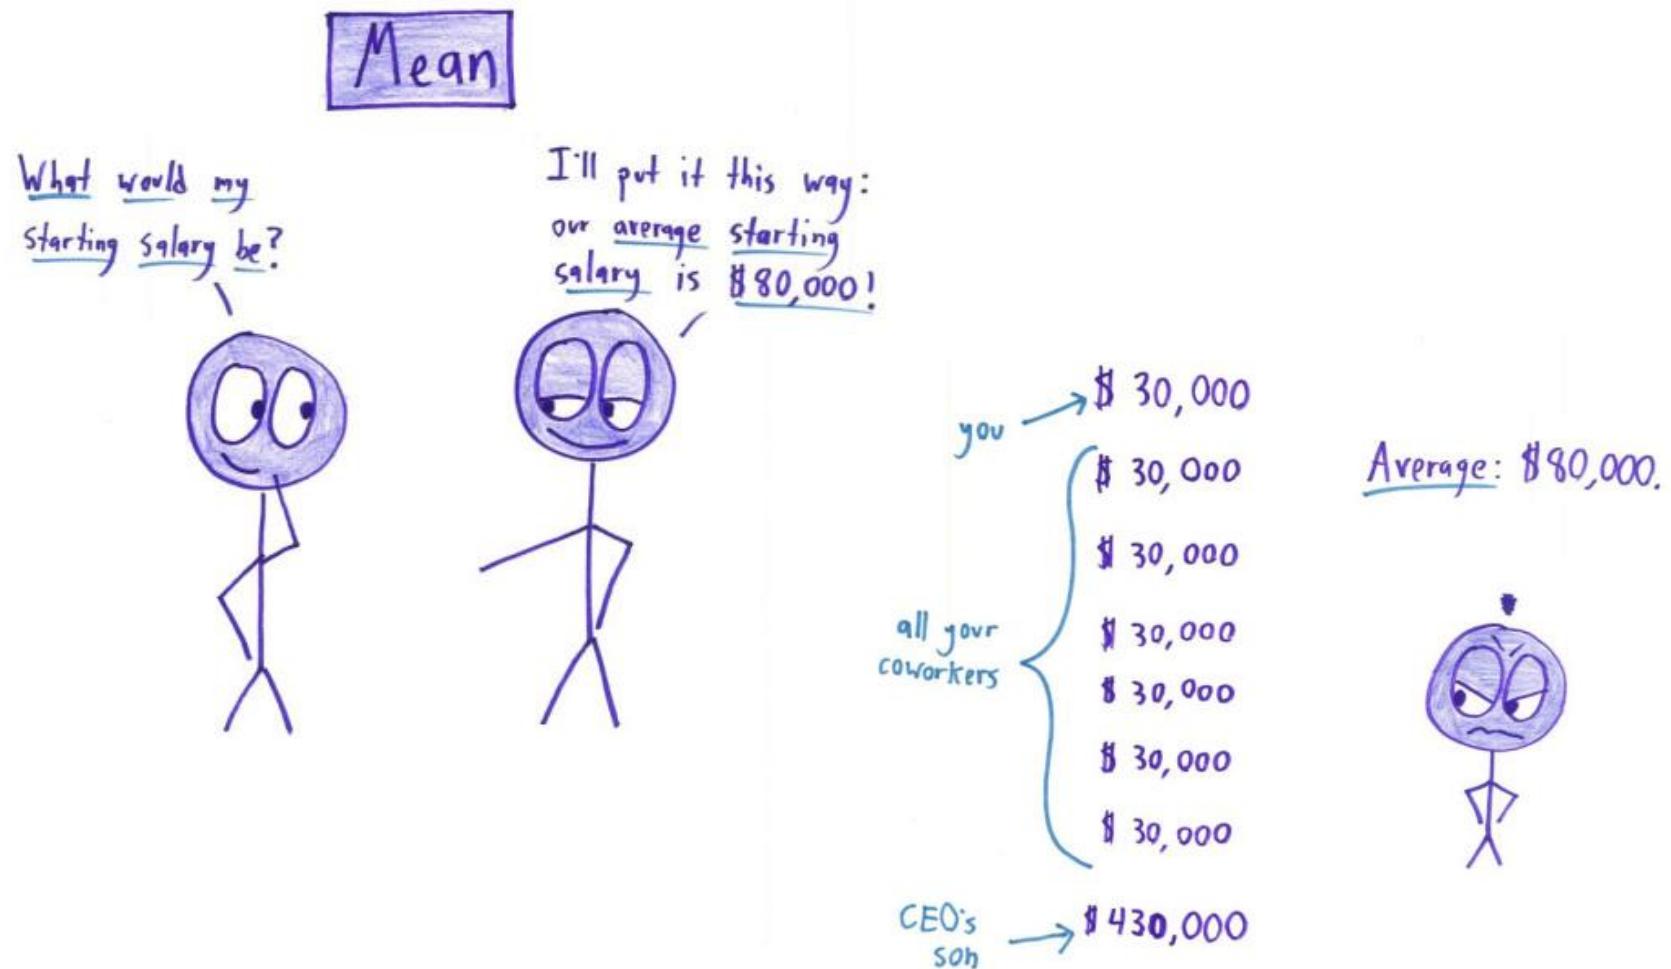


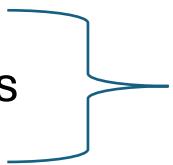
# Lecture 4 – Statistics II



# Today's Learning Outcomes

1. Be able to explain, in plain language, what the standard deviation formula is actually measuring
2. Be able to compile measures of central tendency and variance into a data frame in Python

# Summarizing Data

- Since we want large sample sizes of variable data there's a need to summarize the data (i.e. reduce it to a few easily understood numbers)
- Variance/Uncertainty
  - Standard deviation
  - Confidence intervals
  - Max/Min values

Range of possible values/estimates
- Shape of a distribution
  - Mean
  - Mode
  - Percentiles
    - Median
    - Quartiles
    - Quintiles

Which values are most common relative to other values

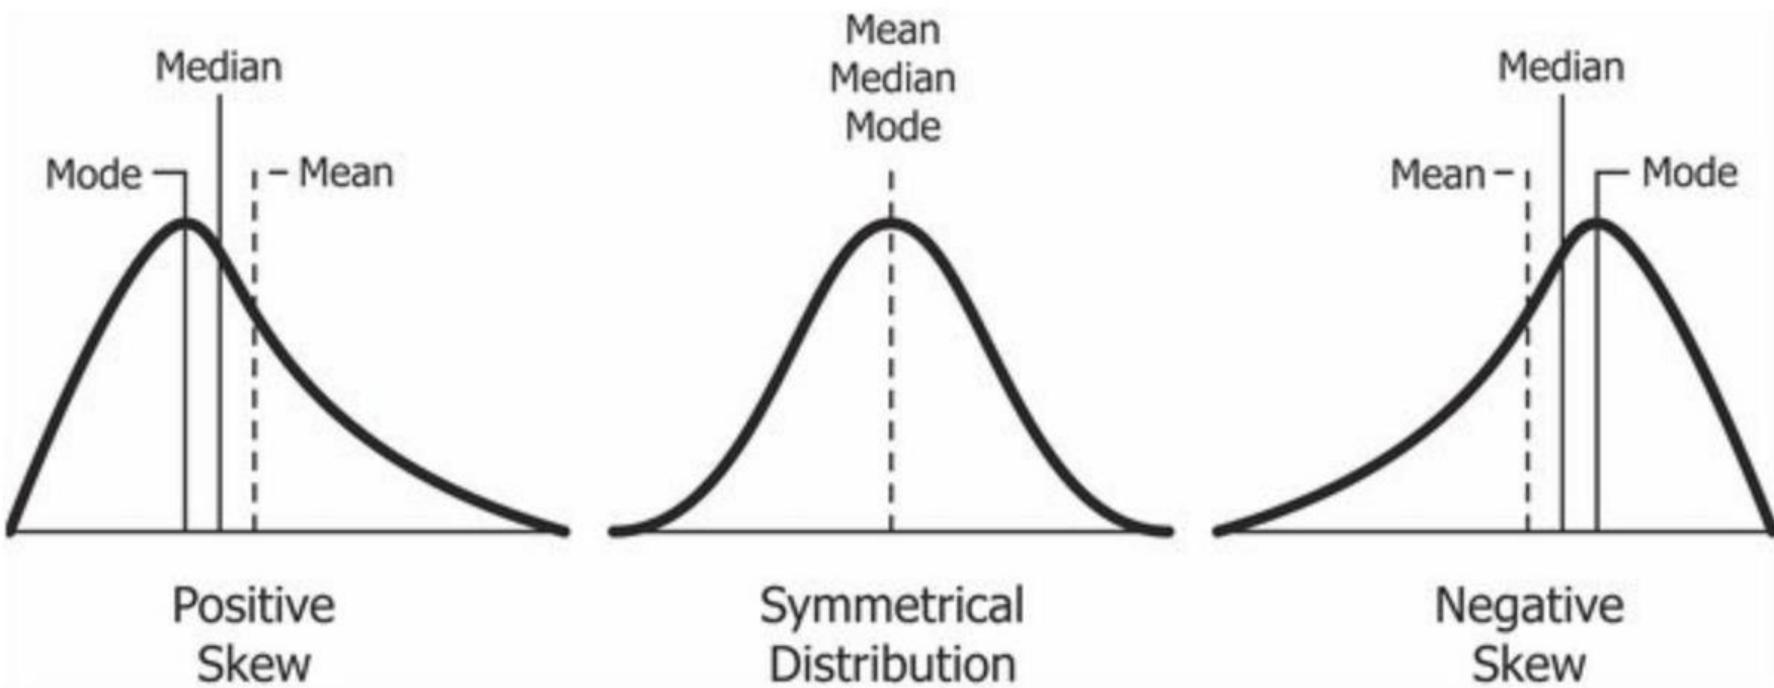
# Shape of Distributions

- Mean, median, mode ← measures of central tendency
  - i.e. where the center of sample values is
- If the goal is to actually find the middle of a sample/distribution the median is the most robust
- mean() and median()
  - No base command for mode (have to make a function)

```
a = [1,2,2,3,3,3]  
# to calculate mean and median, we  
can simply use numpy  
import numpy as np  
a_mean = np.mean(a)  
a_median = np.median(a)
```

```
# to calculate mode, we will need  
to use pandas  
import pandas as pd  
s = pd.Series([1,2,2,3,3])  
print(s.mode())
```

- The mean value is strongly affected by outlier values (i.e. extremes) and mode is vulnerable/uninformative when values are relatively even
  - Mean shifts toward the long tail of a distribution

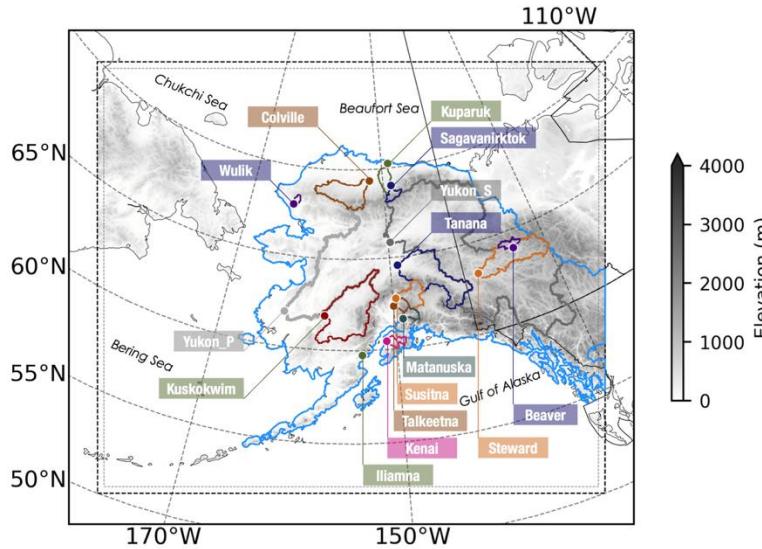


# Case study: how would you report your results?

	RASM	CTSM
Iliamna	<b>0.49</b>	0.32
Wulik	0.17	0.25
Beaver	0.51	0.51
Kuparuk	0.13	0.35
Sagavanirktok	<b>0.62</b>	0.53
Matanuska	0.47	0.59
Talkeetna	<b>0.65</b>	0.55
Kenai	<b>0.50</b>	0.43
Steward	<b>0.71</b>	0.64
Susitna	<b>0.70</b>	0.61
Colville	0.00	0.47
Tanana	<b>0.61</b>	0.56
Kuskokwim	<b>0.22</b>	0.03
Yukon_S	<b>0.58</b>	0.50
Yukon_P	<b>0.60</b>	0.50
median	<b>0.51</b>	0.5
mean	<b>0.464</b>	0.456

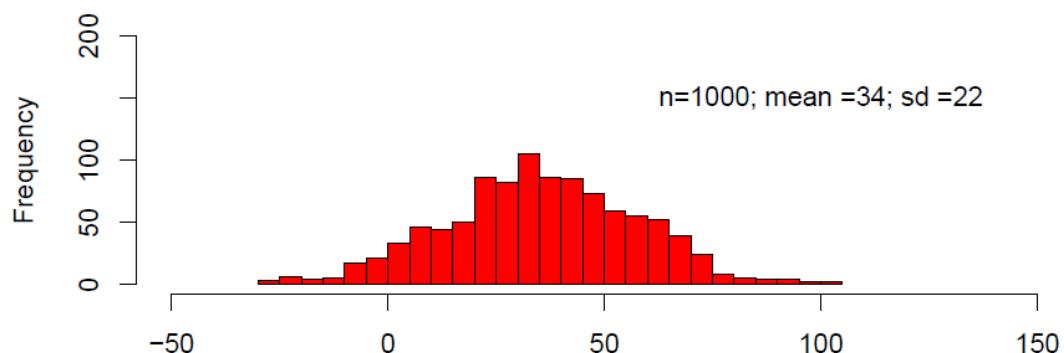
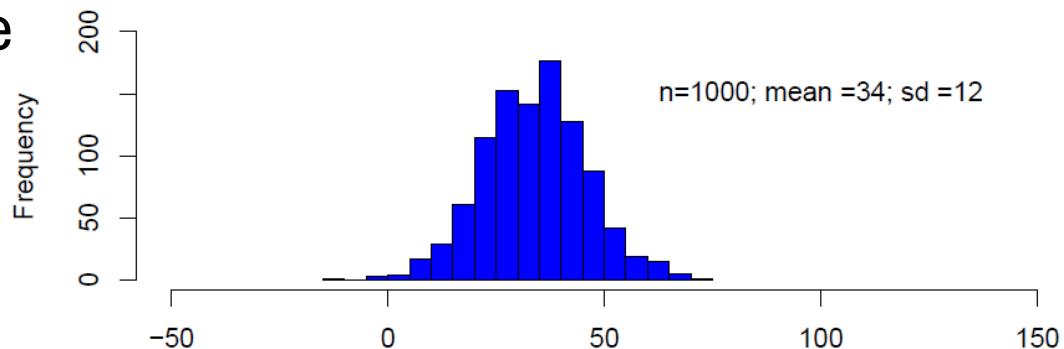
Left table shows the Nash-Sutcliffe Efficiency (**NSE**) of two model simulations(**RASM** versus **CTSM**) against USGS observations for 15 major river basins across.

1. Higher **NSE** values mean better performance
2. The basins are ordered from small basins to large basins (from top to bottom)



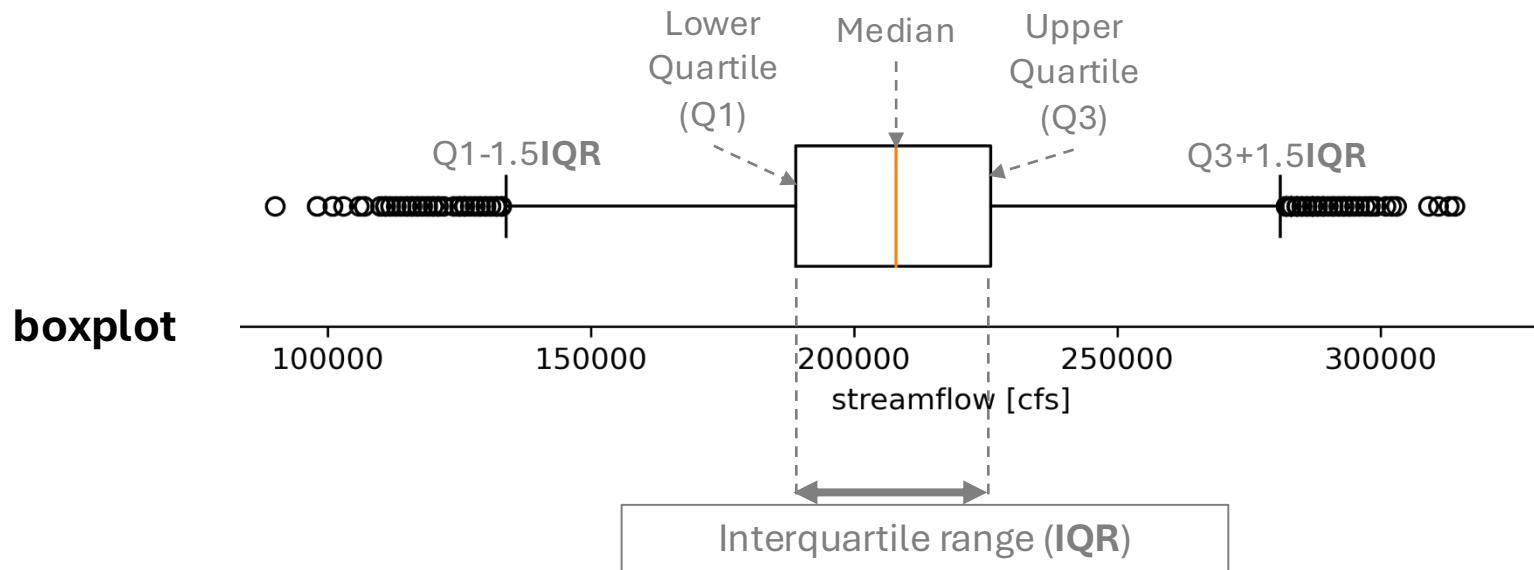
# Measuring Variance of Samples

- Point values alone only tell us so much about data
- We need to measure the spread of values as well
  - Percentiles
  - Standard deviation
  - Max/min



# Percentiles

- Percentiles are the values of a distribution that  $x\%$  of the distribution is less than or equal to
  - The median is the 50<sup>th</sup> percentile
- Common percentiles are 25<sup>th</sup>, 50<sup>th</sup>, and 75<sup>th</sup>
  - Interquartile range of boxplots are the values between the 25<sup>th</sup> and 75<sup>th</sup> percentiles



# Standard Deviation

Python function

```
np.std()
```

$$\sigma = \sqrt{\left(\frac{\sum(x_i - \mu)^2}{N - 1}\right)}$$

$\sigma$  = standard deviation

$x_i$  = a data point value

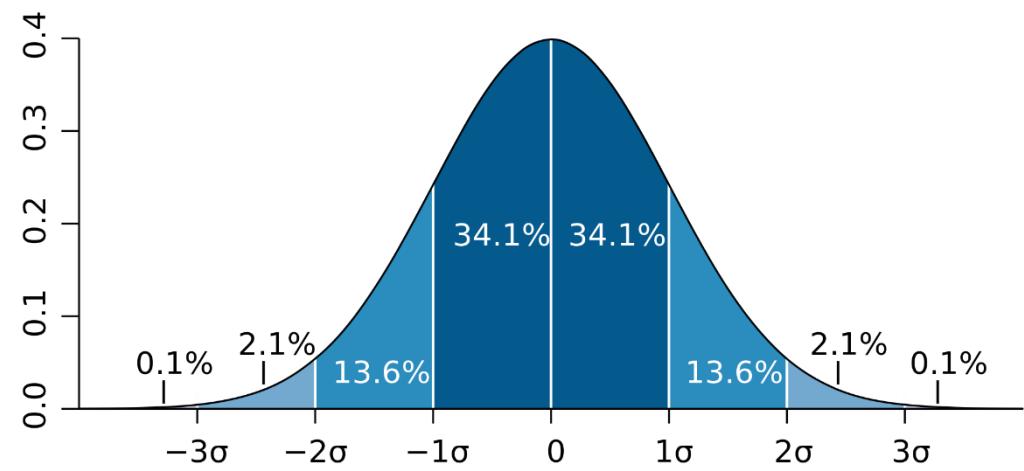
$\mu$  = mean value of the data

N = total number of data points

$\Sigma$  = sum operator

Standard deviation reflects the spread of values assuming the data being sampled is from a **normal distribution**

The 2<sup>nd</sup> power means big deviations from the mean are weighted more heavily (i.e. they indicate a flatter and broader normal distribution with long tails)



Increasing sample size (N) decreases the calculated SD as the chance of values being poor representative of the distribution decreases [i.e. less randomness]

# Summarizing Data

- As with any summary these numbers flatten out variation
- Always useful to plot your actual data and make sure they are appropriate and actually represent what you think they do!
- Open Github Codespace in your homework repo
- Go to “CourseMaterials/coding\_modules/” and double click “lec4\_inclass\_practice\_basic\_statistics.ipynb”

- Four columns of 1000 values, each sampled randomly from a different distribution
1. Load “Lect4\_Data.csv” file into Python

```
import pandas as pd
pd.read_csv()
```
  2. Create a histogram for each of the four datasets on one plot

```
import matplotlib.pyplot as plt
plt.hist()
```
  3. Make a data frame containing the mean, 25<sup>th</sup> percentile, median, 75<sup>th</sup> percentile, and standard deviation of each of the four distributions

```
pd.DataFrame()
```
- If we wanted to get the quintiles (0,20,40,60,80,100) of one of these how would you do so?
    - Just conceptually for now, you don't have to try and write a Python code

# Measuring Uncertainty

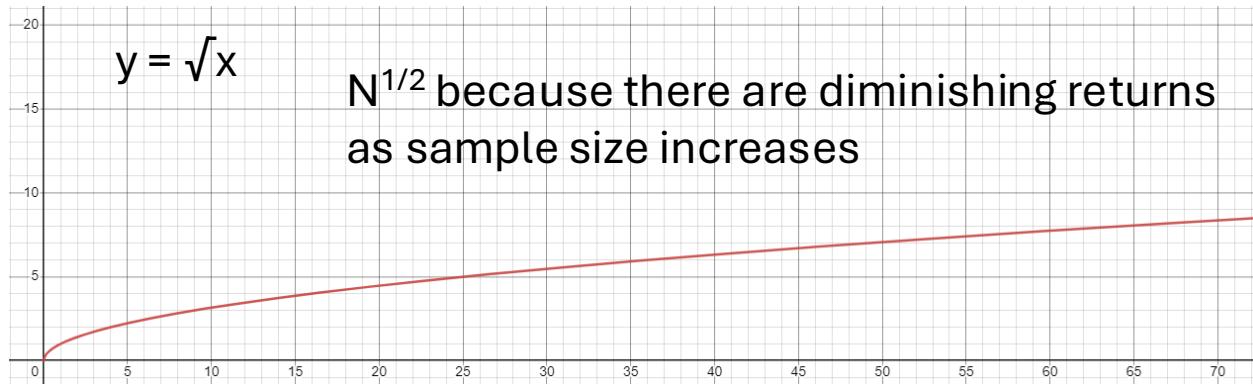
- Acknowledge that we have only a sample estimate (not the population) so if we were to resample we would almost certainly get a slightly different answer
- Confidence intervals (CI) are ways of expressing the uncertainty in an estimated value for a sampling method
  - Can be derived by assuming/fitting a distribution to data and then the range of values expected by that distribution
    - Ex. CI of Correlation coefficient in Python `pearsonr(x1, x2)`
  - Resampling techniques (bootstrap/jackknife/Monte Carlo)

\* Pearsonr needs to be imported from scipy package  
`from scipy.stats import pearsonr`

# Standard Error (SE)

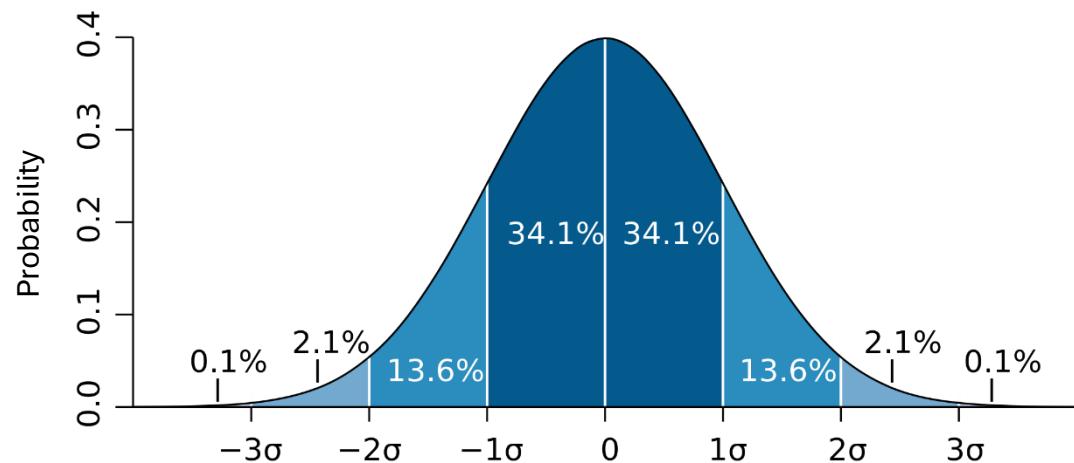
- Standard error (SE) = (standard deviation of the data)/(N)<sup>1/2</sup> ←  $x^{1/2} = \sqrt{x}$
- Measure of uncertainty in the observed mean based on sample size
  - As sample size increases there is **more certainty** in the estimated mean value

Is our estimate likely to change very much with 10,000 versus 20,000 values?



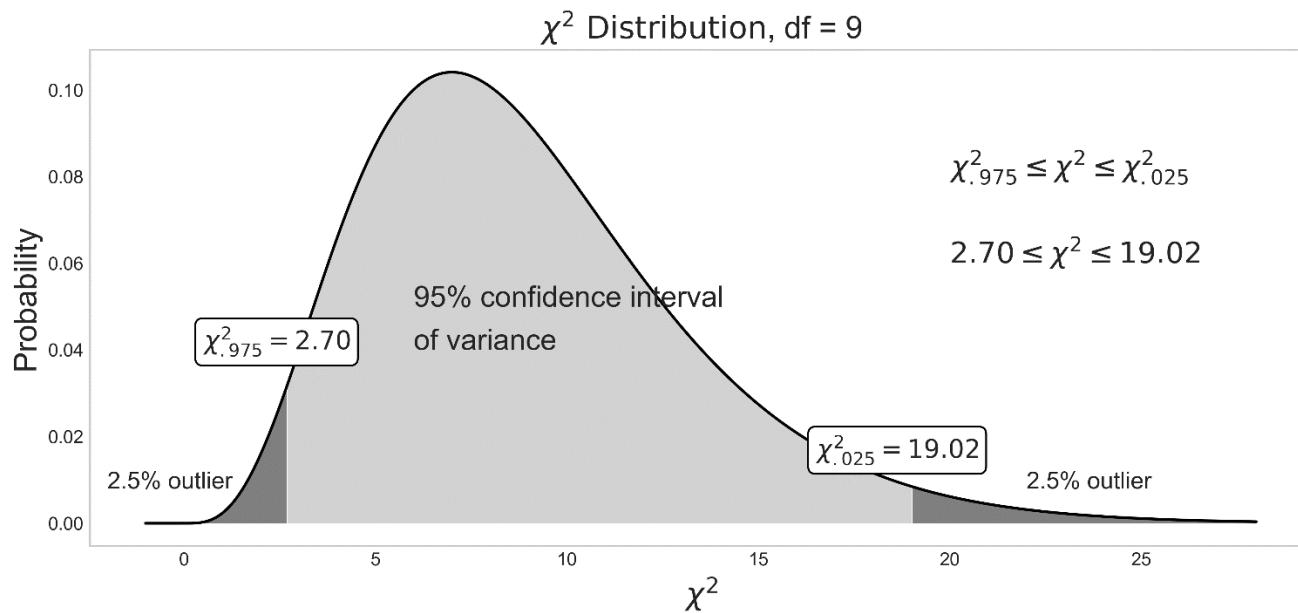
# Uncertainty in Estimates

- Reported confidence intervals are often 95% CI
  - Tied to old practice of using p-values with  $p < 0.05$  being considered statistically significant
- P-value of 0.05  $\approx 5\%$  chance of result by random chance
- $\approx 2$  standard deviations for a normal distribution
  - Chosen totally arbitrarily as the standard



# Uncertainty in Estimates

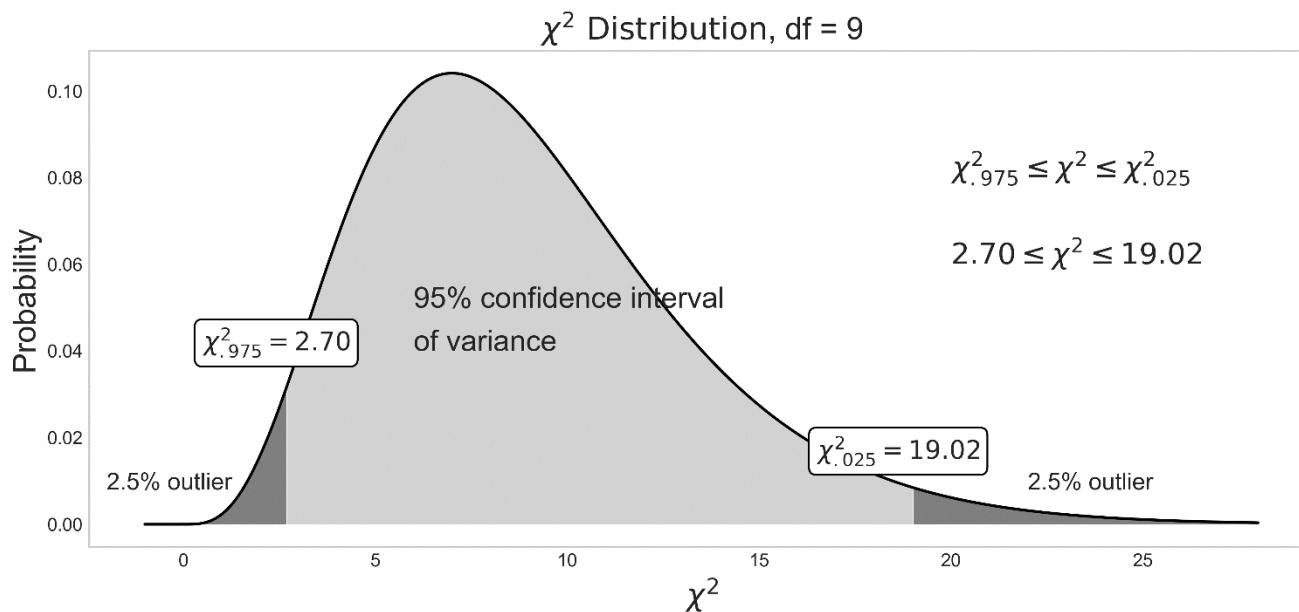
- Lots of data not normal so cannot use 2SD as our standard in many cases (or at least that range is not what most people will think of it as)
  - But 95% of the range of values observed IS useful!
  - Range of values between the 5<sup>th</sup> and 95<sup>th</sup> percentiles



# Uncertainty in Estimates

- Can use same approach as making quintiles to make 95% (or any other range)

```
NinetyFiveCI = [np.percentile(x, 2.5), np.percentile(x, 97.5)]
```



# Today's Learning Outcomes

1. Be able to explain, in plain language, what the standard deviation formula is actually measuring
2. Be able to compile measures of central tendency and variance into a data frame in Python